

Operating System

Chapter 2: Processes and Threads

Prepared By:

Amit K. Shrivastava

Asst. Professor

Nepal College Of Information Technology

Process Concept

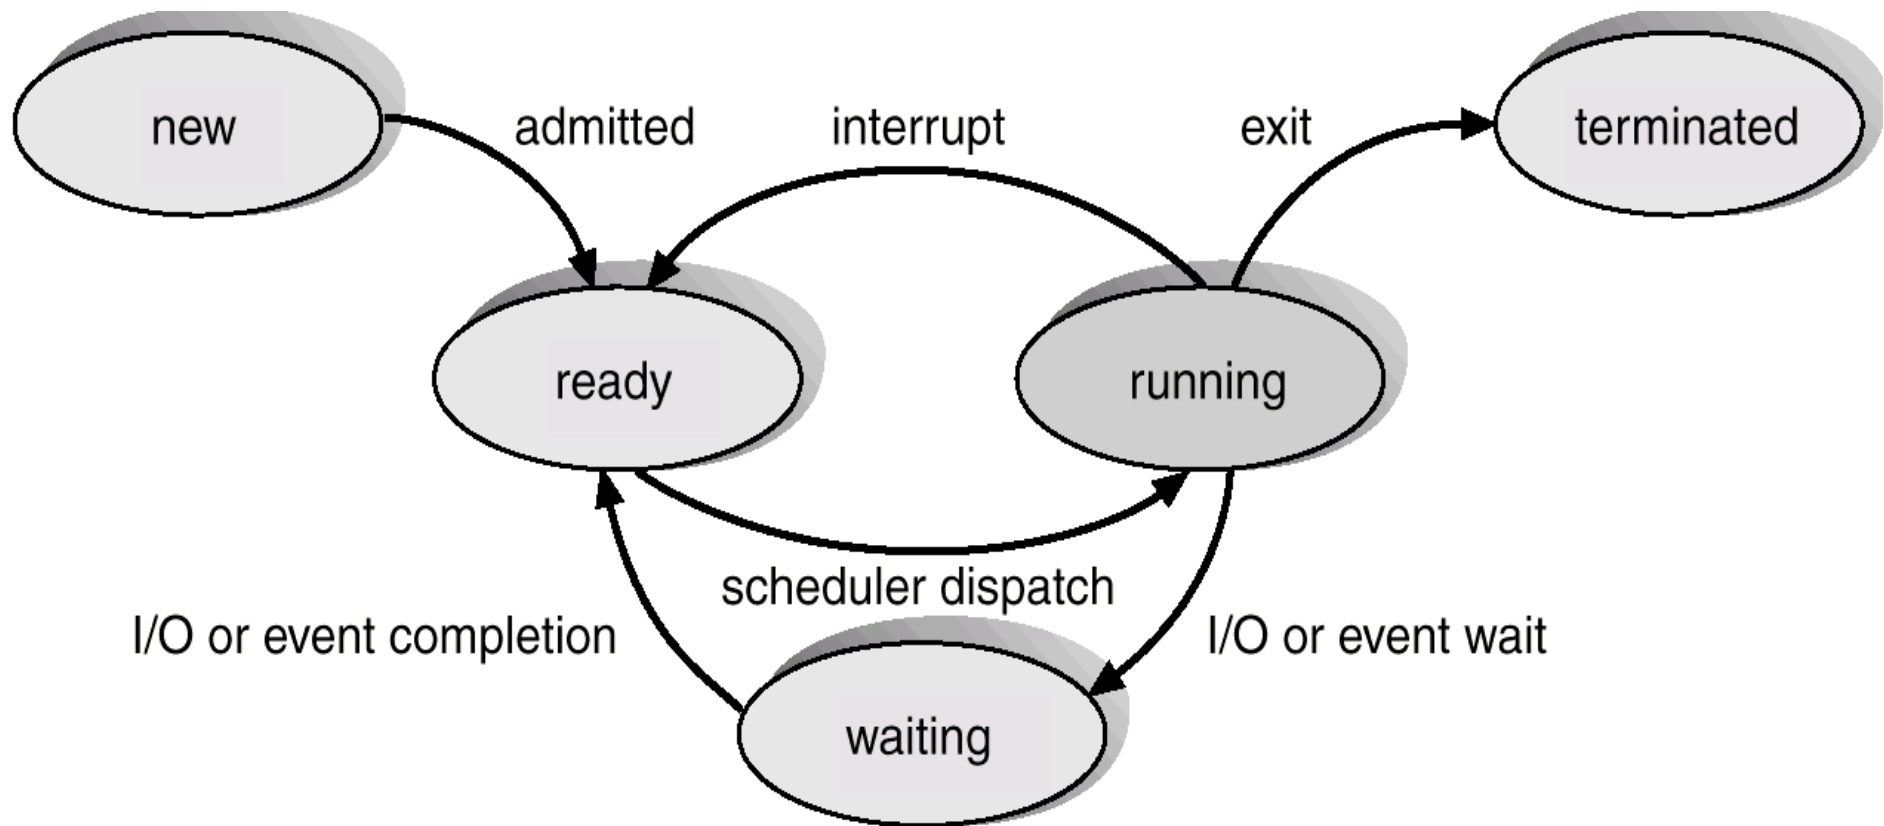
- An operating system executes a variety of programs:
 - ✦ Batch system – jobs
 - ✦ Time-shared systems – user programs or tasks
- The terms job and process are used almost interchangeably.
- Process – a program in execution; process execution must progress in sequential fashion.
- A process includes:
 - ✦ program counter
 - ✦ stack
 - ✦ data section

Process State

■ As a process executes, it changes state

- ✦ new: The process is being created.
- ✦ running: Instructions are being executed.
- ✦ waiting: The process is waiting for some event to occur.
- ✦ ready: The process is waiting to be assigned to a process.
- ✦ terminated: The process has finished execution.

Diagram of Process State



Process Control Block (PCB)

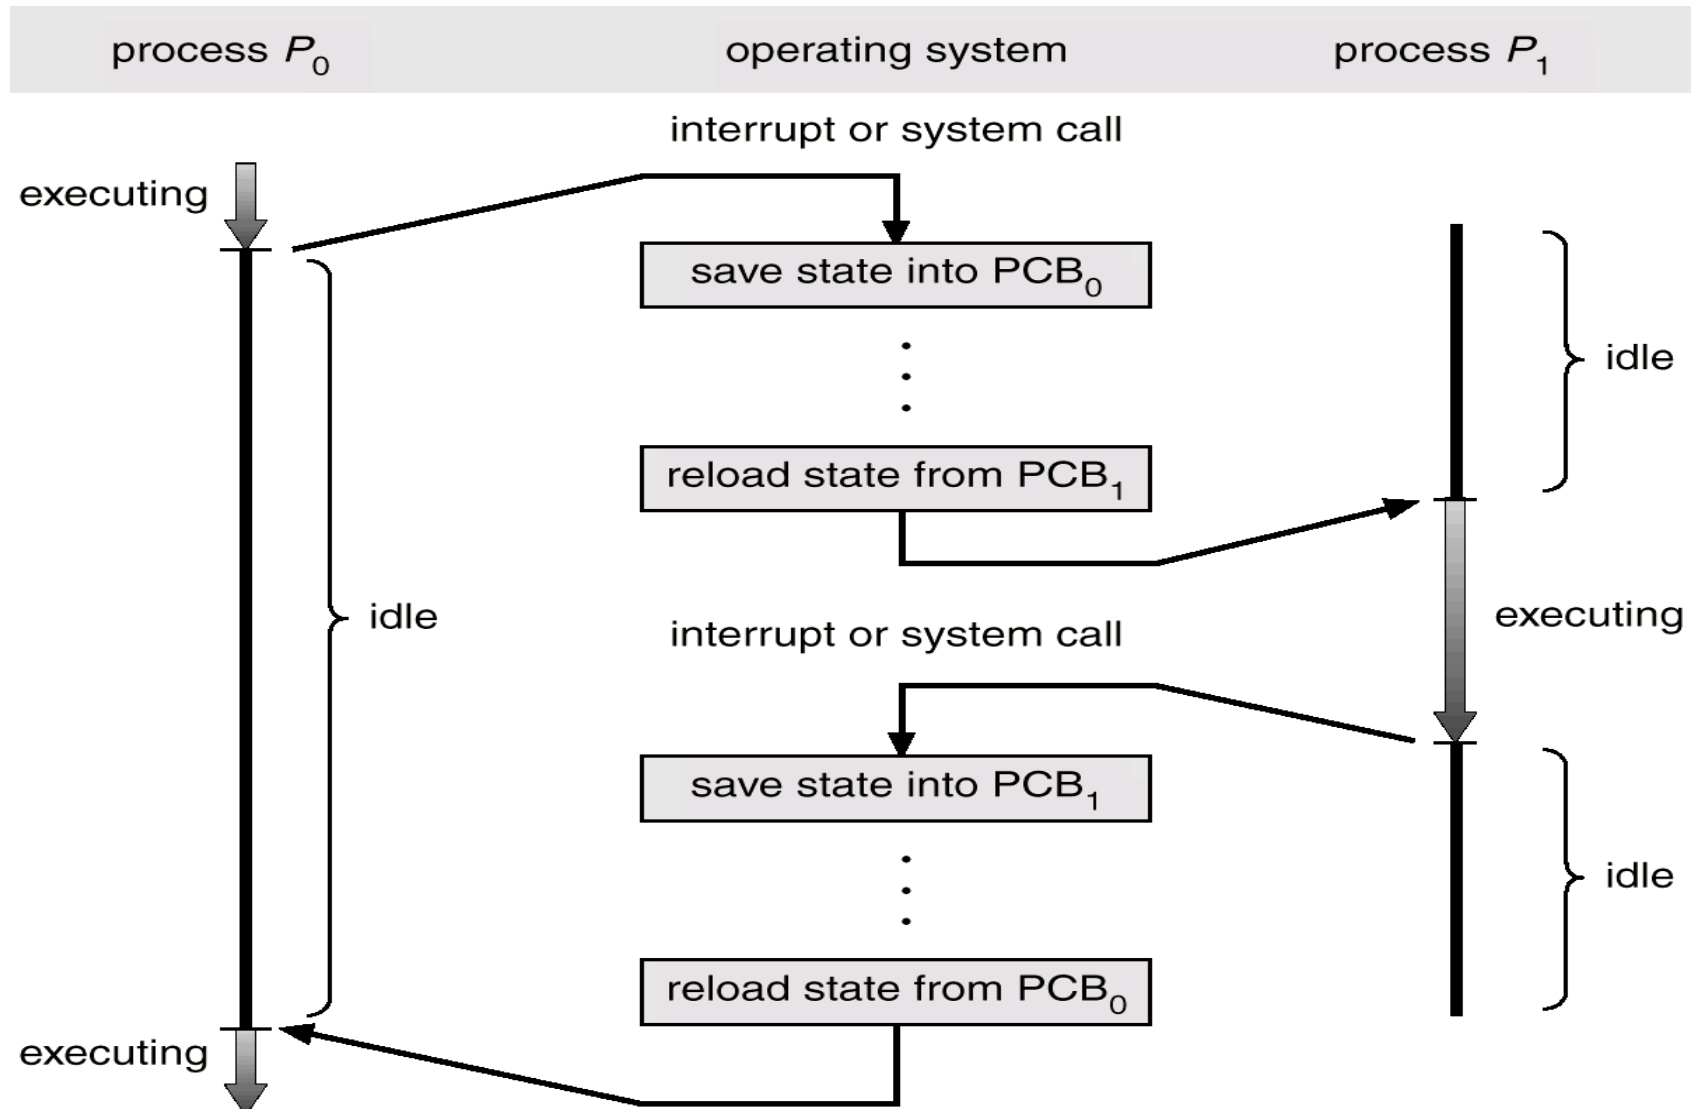
Information associated with each process.

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information

Process Control Block (PCB)

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
• • •	

CPU Switch From Process to Process



Cooperating Processes

- Independent process cannot affect or be affected by the execution of another process.
- Cooperating process can affect or be affected by the execution of another process

Advantages of process cooperation:

- Information sharing: Since several users may be interested in the same piece of information (for instance, a shared file), we must provide an environment to allow concurrent access to these types of resources.
- Computation speedup: If we want a particular task to run faster, we must break it into subtasks, each of which will be executing in parallel with the others. Such a speedup can be achieved only if the computer has multiple processing elements (such as CPUs or I/O channels).

-Modularity: We may want to construct the system in a modular fashion, dividing the system functions into separate processes or threads,

-Convenience: Even an individual user may have many tasks on which to work at one time. For instance, a user may be editing, printing, and compiling in parallel.

IPC(Inter-Process Communication):

- Processes frequently need to communicate with other processes i.e they need to exchange data and information. Thus there is need a need for communication between processes, preferably in a well-structured way not using interrupts. This communication mechanism between processes are known as Inter-Process Communication.
- There are two fundamental models of interprocess communication: shared memory and message passing.
- In the shared-memory model, a region of memory that is shared by cooperating processes is established. Processes can then exchange information by reading and writing data to the shared region.
- In the message-passing model, communication takes place by means of messages exchanged between the cooperating processes.

IPC(Inter-Process Communication) contd....:

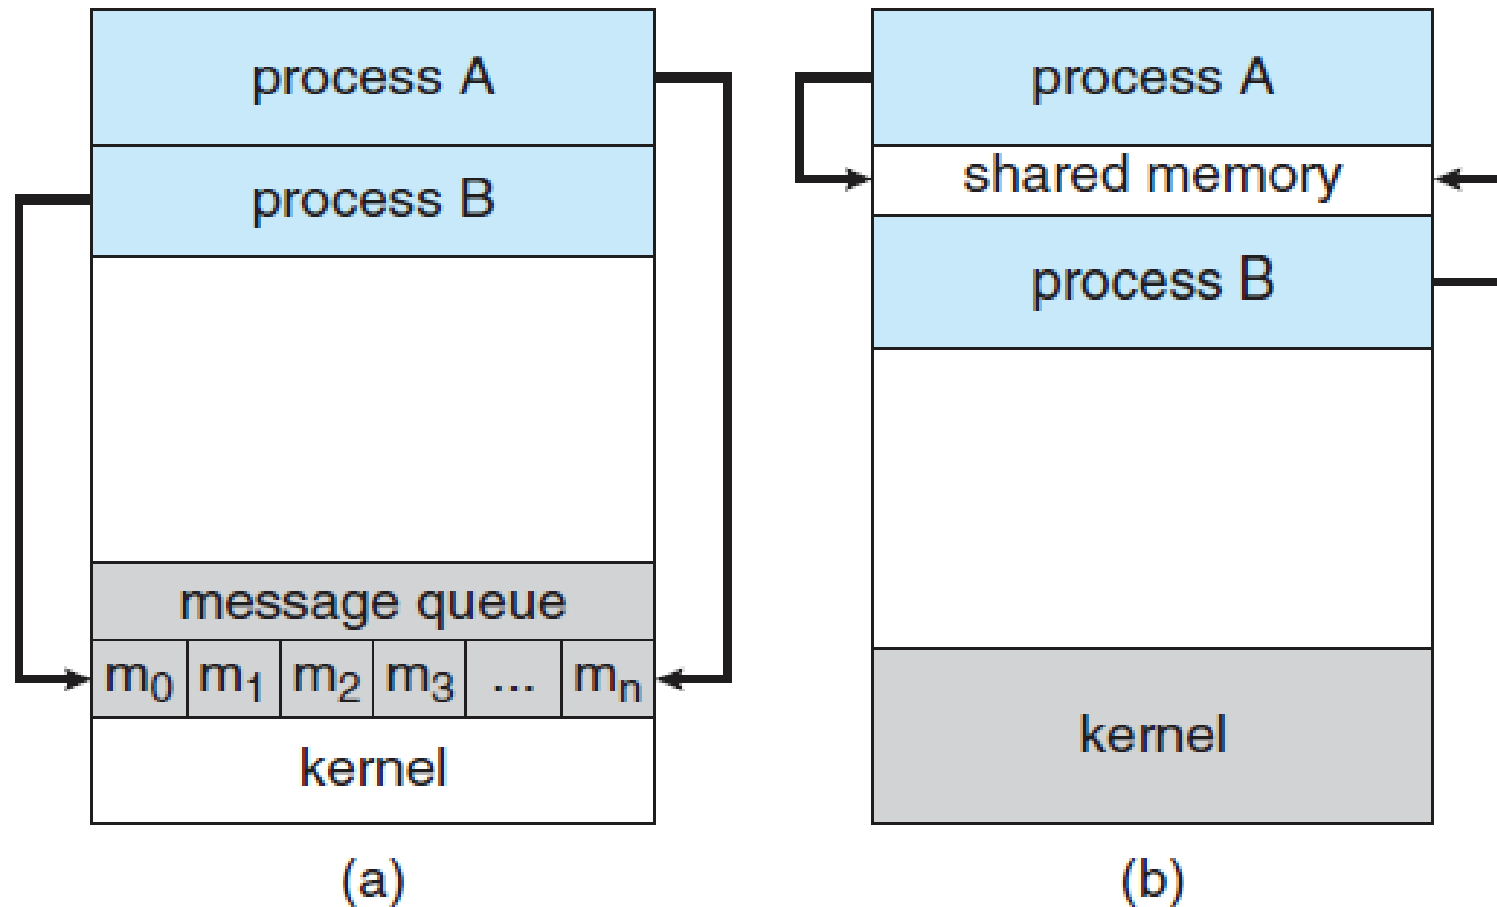


Figure 3.12 Communications models. (a) Message passing. (b) Shared memory.

Race Condition:

- In some operating systems, processes that are working together may share some common storage that each one can read and write.
- Situations like this, where two or more processes are reading or writing some shared data and the final result depends on who runs precisely when, are called **race conditions**.

Critical Regions:

- Sometimes a process has to access shared memory or files, or do other critical things that can lead to races. That part of the program where the shared memory is accessed is called the **critical region** or **critical section**.
- If we could arrange matters such that no two processes were ever in their critical regions at the same time, we could avoid race conditions.

Critical Regions(contd..):

- Although this requirement avoids race conditions, this is not sufficient for having parallel processes cooperate correctly and efficiently using shared data. We need four conditions to hold to have a good solution:
 1. No two processes may be simultaneously inside their critical regions.
 2. No assumptions may be made about speeds or number of CPUs.
 3. No process running outside its critical region may block other processes.
 4. No process should have to wait forever to enter its critical region.
- The behavior that we want is shown in Fig. of next slide.

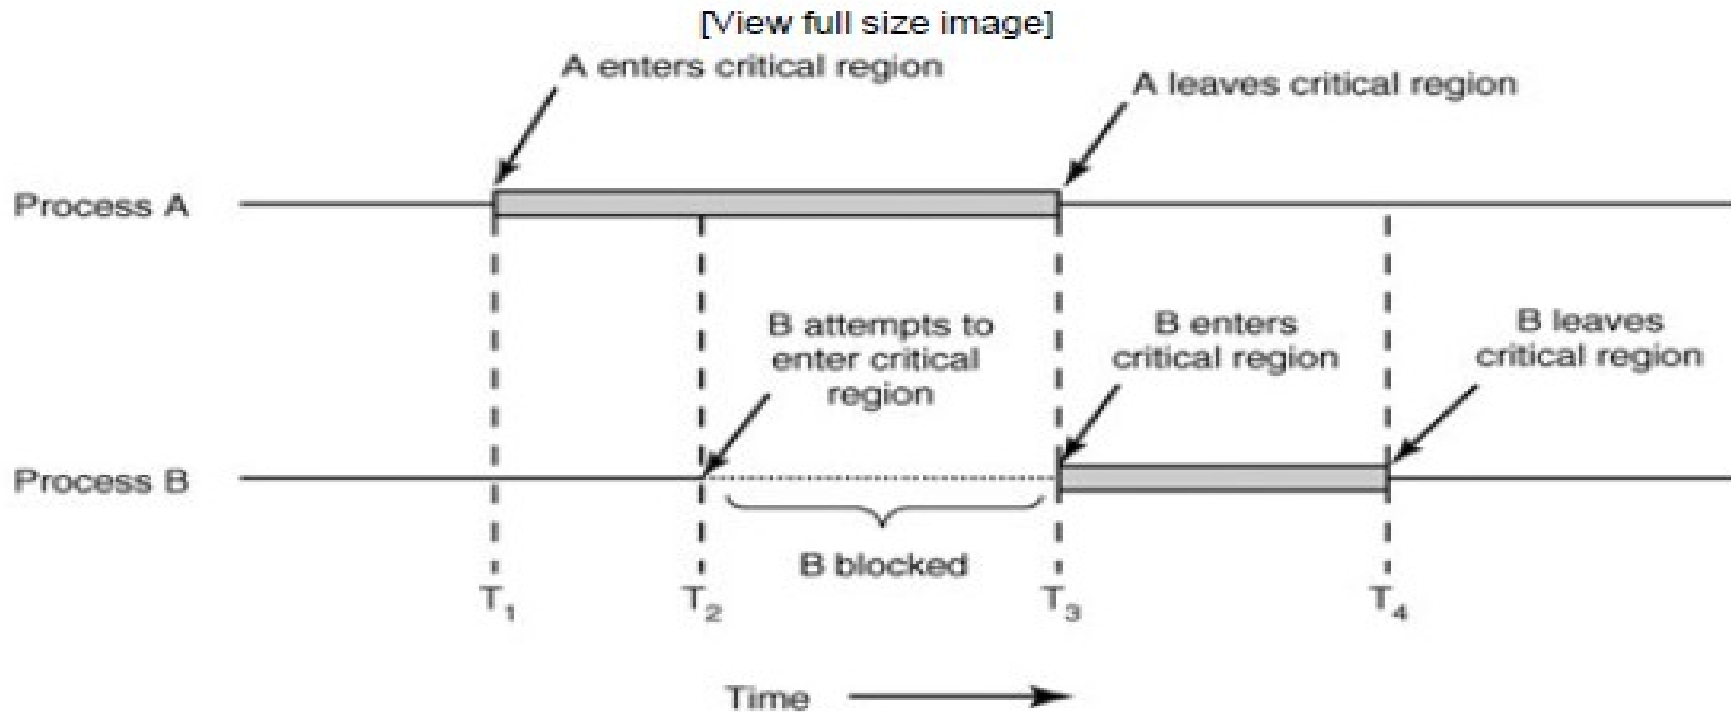


Fig: Mutual Exclusion using Critical Regions

Here process A enters its critical region at time T_1 . A little later, at time T_2 process B attempts to enter its critical region but fails because another process is already in its critical region and we allow only one at a time. Consequently, B is temporarily suspended until time T_3 when A leaves its critical region, allowing B to enter immediately. Eventually B leaves (at T_4) and we are back to the original situation with no processes in their critical regions.

Mutual Exclusion:

- How do we avoid race conditions? The key to preventing trouble here and in many other situations involving shared memory, shared files, and shared everything else is to find some way to prohibit more than one process from reading and writing the shared data at the same time. i.e, what we need is mutual exclusion some way of making sure that if one process is using a shared variable or file, the other processes will be excluded from doing the same thing.

Peterson's Algorithm:

- In 1981, G.L P discovered a much simpler way to achieve mutual exclusion

```
#define FALSE 0
#define TRUE  1
#define N      2          /* number of processes */

int turn;                 /* whose turn is it? */
int interested[N];        /* all values initially 0 (FALSE)*/
void enter_region(int process) /* process is 0 or 1 */
{
    int other;             /* number of the other process */
    other = 1 - process;   /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;        /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */;
}
void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```


Peterson's Algorithm Explanation

- Before using the shared variables (i.e., before entering its critical region), each process calls `enter_region` with its own process number, 0 or 1, as the parameter. This call will cause it to wait, if need be, until it is safe to enter. After it has finished with the shared variables, the process calls `leave_region` to indicate that it is done and to allow the other process to enter, if it so desires.
- Let us see how this solution works. Initially, neither process is in its critical region. Now process 0 calls `enter_region`. It indicates its interest by setting its array element and sets `turn` to 0. Since process 1 is not interested, `enter_region` returns immediately. If process 1 now calls `enter_region`, it will hang there until `interested[0]` goes to `FALSE`, an event that only happens when process 0 calls `leave_region` to exit the critical region.
- Now consider the case that both processes call `enter_region` almost simultaneously. Both will store their process number in `turn`. Whichever store is done last is the one that counts; the first one is lost. Suppose that process 1 stores last, so `turn` is 1. When both processes come to the while statement, process 0 executes it zero times and enters its critical region. Process 1 loops and does not enter its critical region.

TSL(Test and Set Lock)

➤ TSL is a hardware solution to the mutual exclusion problem. The key to success here is to have a single hardware instruction that reads a variable, stores its value in a save area, and sets the variable to a certain value. This instruction, often called test and set, once initiated will complete all of these functions without interruption. The instruction is like

TSL REGISTER, LOCK

➤ (Test and Set Lock) that works as follows. It reads the contents of the memory word lock into register RX and then stores a nonzero value at the memory address lock. The operations of reading the word and storing into it are guaranteed to be indivisible—no other processor can access the memory word until the instruction is finished. The CPU executing the TSL instruction locks the memory bus to prohibit other CPUs from accessing memory until it is done.

➤ To use the TSL instruction, we will use a shared variable, lock, to coordinate access to shared memory. When lock is 0, any process may set it to 1 using the TSL instruction and then read or write the shared memory. When it is done, the process sets lock back to 0 using an ordinary move instruction.

TSL(contd..)

do {

acquire lock

Critical section

Release lock

remainder section

} while(TRUE);

fig: Solution to the critical-section problem using locks

enter_region:

TSL REGISTER,LOCK	/ copy lock to register and set lock to 1
CMP REGISTER,#0	/ was lock zero?
JNE enter_region	/ if it was nonzero, lock was set, so loop
RET	/ return to caller; critical region entered

leave_region:

MOVE LOCK,#0	/ store a 0 in lock
RET	/ return to caller

Figure : Entering and leaving a critical region using the TSL instruction.

Producer and Consumer Problem

- ◆ To illustrate the concept of cooperating processes, let us consider the producer-consumer problem, which is a common paradigm for cooperating processes.
- ◆ A producer process produces information that is consumed by a consumer process. For example, a print program produces characters that are consumed by the printer driver. A compiler may produce assembly code, which is consumed by an assembler.
- ◆ To allow producer and consumer processes to run concurrently, we must have available a buffer of items that can be filled by the producer and emptied by the consumer.
- ◆ A producer can produce one item while the consumer is consuming another item.
- ◆ The producer and consumer must be synchronized, so that the consumer does not try to consume an item that has not yet been produced. In this situation, the consumer must wait until an item is produced.

◆ Unbounded-buffer places no practical limit on the size of the buffer- The consumer may have to wait for new items, but the producer can always produce new items.

◆ Bounded-buffer assumes that there is a fixed buffer size. In this case, the consumer must wait if the buffer is empty, and the producer must wait if the buffer is full.

◆ The buffer may either be provided by the operating system through the use of an interprocess-communication (IPC) facility or by explicitly coded by the application programmer with the use of shared memory.

- Trouble arises when the producer wants to put a new item in the buffer, but it is already full. The solution is for the producer to go to sleep, to be awakened when the consumer has removed one or more items. Similarly, if the consumer wants to remove an item from the buffer and sees that the buffer is empty, it goes to sleep until the producer puts something in the buffer and wakes it up.

```

#define N 100                                     /* number of slots in the buffer */
int count = 0;                                   /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                               /* repeat forever */
        item = produce_item();                  /* generate next item */
        if (count == N) sleep();                /* if buffer is full, go to sleep */
        insert_item(item);                      /* put item in buffer */
        count = count + 1;                      /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);      /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                               /* repeat forever */
        if (count == 0) sleep();                /* if buffer is empty, got to sleep */
        item = remove_item();                  /* take item out of buffer */
        count = count - 1;                      /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer);  /* was buffer full? */
        consume_item(item);                    /* print item */
    }
}

```

Figure 2-27. The producer-consumer problem with a fatal race condition.

Monitors

- High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.
- The representation of a monitor type consists of declarations of variables whose values define the state of an instance of the type, as well as the bodies of procedures or functions that implement operations on the type.

```
monitor monitor-name
{
    shared variable declarations
    procedure body P1 (...) {
        ...
    }
    procedure body P2 (...) {
        ...
    }
    procedure body Pn (...) {
        ...
    }
    {
        initialization code
    }
}
```

Monitors(contd..)

◆ To allow a process to wait within the monitor, a **condition** variable must be declared, as

condition x, y;

◆ Condition variable can only be used with the operations **wait** and **signal**.

The operation

x.wait();

means that the process invoking this operation is suspended until another process invokes

x.signal();

The **x.signal** operation resumes exactly one suspended process. If no process is suspended, then the signal operation has no effect.




Figure 7.20 Schematic view of a monitor.

Message Passing

When Process interact with one another,two fundamental requirements must be satisfied: synchronization and communication. One approach to providing both of these functions is message passing, i.e message passing allows information exchange between machines. This method of interprocess communication uses two primitives, send and receive, system calls. As such, they can easily be put into library procedures, such as

 Send(destination, & message);

and

 Receive(source, & message);

The former call sends a message to a given destination and the latter one receives a message from a given source.

Design Issues for Message Passing Systems

- Message passing systems have many challenging problems and design issues, especially if the communicating processes are on different machines connected by a network. For example, messages can be lost by the network. To guard against lost messages, the sender and receiver can agree that as soon as a message has been received, the receiver will send back a special acknowledgement message. If the sender has not received the acknowledgement within a certain time interval, it retransmits the message.
- To avoid duplication, consecutive sequence numbers can be put in each original messages.

Classical IPC Problems

There are number of different synchronization problem. Some of them are:

The Dining-Philosophers Problem:

Statement: Consider five philosophers who spend their lives thinking and eating. The philosophers share a common circular table surrounded by five chairs, each belonging to one philosopher. In the centre of the table there is a bowl of rice, and the table is laid with five single chopsticks. When a philosophers gets hungry she tries to pick up her left and right chopstick, one at a time, in either order. If successful in acquiring two chopsticks, she eats for a while, then puts down the chopsticks and continues to think. The key question is: Can you write a program for each philosopher that does what it is supposed to do and never gets stuck?



fig:The situation of the dining philosophers

Solution to the Dining-Philosophers Problem:

One simple solution is to represent each chopstick with a semaphore. A philosopher tries to grab a chopstick by executing a wait() operation on that semaphore; she releases her chopsticks by executing the signal() operation on the appropriate semaphores. Thus , the shared data are

semaphore chopstick [5];

The structure of Philosopher i is shown below:

```
do {  
    wait ( chopstick[i] );  
    wait ( chopStick[ (i + 1) % 5] );  
    . . . . .  
    // eat  
    . . . . .  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
    . . .  
    // think  
    . . .  
} while (true) ;
```

fig: the structure of philosopher i.

But if this solution simply applied deadlock will occur.

Solution to the Dining-Philosophers Problem(contd...):

Several possible remedies to the deadlock problem are listed. Here is a solution to the dining-philosophers problem that ensures freedom from deadlocks.

- Allow at most four philosophers to be sitting simultaneously at the table.
- Allow a philosopher to pick up her chopsticks only if both chopsticks are available (to do this she must pick them up in a critical section).
- Use an asymmetric solution; that is, an odd philosopher picks up first her left chopstick and then her right chopstick, whereas an even philosopher picks up her right chopstick and then her left chopstick.

Readers Writer Problems:

- The dining philosophers problem is useful for modeling processes that are competing for exclusive access to a limited number of resources, such as I/O devices. Another famous problem is the readers and writers problem which models access to a database (Courtois et al., 1971). Imagine, for example, an airline reservation system, with many competing processes wishing to read and write it. It is acceptable to have multiple processes reading the database at the same time, but if one process is updating (writing) the database, no other process may have access to the database, not even a reader. The question is how do you program the readers and the writers? One solution is shown below.

```
typedef int semaphore; /* use your imagination */
semaphore mutex = 1; /* controls access to 'rc' */
semaphore db = 1; /* controls access to the database */
int rc = 0; /* # of processes reading or wanting to */
```

Readers Writer problems(contd..)

```
void reader(void)
{
    while (TRUE){ /* repeat forever */
        down(&mutex); /* get exclusive access to 'rc' */
        rc = rc + 1; /* one reader more now */
        if (rc == 1) down(&db); /* if this is the first reader ... */
        up(&mutex); /* release exclusive access to 'rc' */
        read_data_base(); /* access the data */
        down(&mutex); /* get exclusive access to 'rc' */
        rc = rc - 1; /* one reader fewer now */
        if (rc == 0) up(&db); /* if this is the last reader ... */
        up(&mutex); /* release exclusive access to 'rc' */
        use_data_read(); /* noncritical region */
    }
}

void writer(void)
{
    while (TRUE){ /* repeat forever */
        think_up_data(); /* noncritical region */
        down(&db); /* get exclusive access */
        write_data_base(); /* update the data */
        up(&db); /* release exclusive access */
    }
}
```

fig: A solution to the readers writer problem

Readers Writer problems(contd..)

- In this solution, the first reader to get access to the data base does a down on the semaphore *db* . Subsequent readers merely have to increment a counter, *rc* . As *readers leave, they decrement* the counter and the last one out does an up on the semaphore, allowing a blocked writer, if there is one, to get in.
- Suppose that while a reader is using the data base, another reader comes along. Since having two readers at the same time is not a problem, the second reader is admitted. A third and subsequent readers can also be admitted if they come along.
- Now suppose that a writer comes along. The writer cannot be admitted to the data base, since writers must have exclusive access, so the writer is suspended. Later, additional readers show up. As long as at least one reader is still active, subsequent readers are admitted. As a consequence of this strategy, as long as there is a steady supply of readers, they will all get in as soon as they arrive. The writer will be kept suspended until no reader is present. If a new reader arrives, say, every 2 seconds, and each reader takes 5 seconds to do its work, the writer will never get in.
- To prevent this situation, the program could be written slightly differently: When a reader arrives and a writer is waiting, the reader is suspended behind the writer instead of being admitted immediately. In this way, a writer has to wait for readers that were active when it arrived to finish but does not have to wait for readers that came along after it. The disadvantage of this solution is that it achieves less concurrency and thus lower performance.

The Sleeping Barber Problem

- Another classical IPC problem takes place in a barber shop. The barber shop has one barber, one barber chair, and n chairs for waiting customers, if any, to sit on. If there are no customers present, the barber sits down in the barber chair and falls asleep, as illustrated in [Figure 2-35]. When a customer arrives, he has to wake up the sleeping barber. If additional customers arrive while the barber is cutting a customer's hair, they either sit down (if there are empty chairs) or leave the shop (if all chairs are full). The problem is to program the barber and the customers without getting into race conditions. This problem is similar to various queueing situations, such as a multiperson helpdesk with a computerized call waiting system for holding a limited number of incoming calls.
- This solution uses three semaphores, *customers*, which counts waiting customers (excluding the customer in the barber chair, who is not waiting), *barbers*, the number of barbers (0 or 1) who are idle, waiting for customers, and *mutex*, which is used for mutual exclusion. We also need a variable, *waiting*, which also counts the waiting customers. The reason for having *waiting* is that there is no way to read the current value of a semaphore. In this solution, a customer entering the shop has to count the number of waiting customers. If it is less than the number of chairs, he stays; otherwise, he leaves.
- Our solution is shown [below]. When the barber shows up for work in the morning, he executes the procedure *barber*, causing him to block on the semaphore *customers* because it is initially 0. The barber then goes to sleep, as shown in [Figure 2.35]. He stays asleep until the first customer shows up.
- When a customer arrives, he executes *customer*, starting by acquiring *mutex* to enter a critical region. If another customer enters shortly thereafter, the second one will not be able to do anything until the first one has released *mutex*. The customer then checks to see if the number of waiting customers is less than the number of chairs. If not, he releases *mutex* and leaves without a haircut.

The sleeping Barber Problem(contd...)

- If there is an available chair, the customer increments the integer variable, *waiting*. Then he does an *Up on the semaphore customers*, thus waking up the barber. At this point, the customer and the barber are both awake. When the customer releases *mutex*, the barber grabs it, does some housekeeping, and begins the haircut.
- When the haircut is over, the customer exits the procedure and leaves the shop. Unlike our earlier examples, there is no loop for the customer because each one gets only one haircut. The barber loops, however, to try to get the next customer. If one is present, a haircut is given. If not, the barber goes to sleep.



```

#define CHAIRS 5 /* # chairs for waiting customers */
typedef int semaphore; /* use your imagination */
semaphore customers = 0; /* # of customers waiting for service */
semaphore barbers = 0; /* # of barbers waiting for customers */
semaphore mutex = 1; /* for mutual exclusion */
int waiting = 0; /* customer are waiting (not being cut) */
void barber(void)
{
while (TRUE) {
down(&customers); /* go to sleep if # of customers is 0 */
down(&mutex); /* acquire access to 'waiting' */
waiting = waiting - 1; /* decrement count of waiting customers */
up(&barbers); /* one barber is now ready to cut hair */
up(&mutex); /* release 'waiting' */
cut_hair(); /* cut hair (outside critical region */
}
}
void customer(void)
{
down(&mutex); /* enter critical region */
if (waiting < CHAIRS) { /* if there are no free chairs, leave */
waiting = waiting + 1; /* increment count of waiting customers */
up(&customers); /* wake up barber if necessary */
up(&mutex); /* release access to 'waiting' */
down(&barbers); /* go to sleep if # of free barbers is 0 */
get_haircut(); /* be seated and be served */
} else {
up(&mutex); /* shop is full; do not wait */
}
}

```

fig:A solution to the Sleeping Barber Problem

Deadlocks

Introduction:

➤ In a multiprogramming environment, several processes may compete for a finite number of resources. A process requests resources; if the resources are not available at that time, the process enters a wait state. Waiting processes may never again change state, because the resources they have requested are held by other waiting processes. This situation is called a deadlock

OR

➤ A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.

➤ Example

- ◆ System has 2 tape drives.
- ◆ P1 and P2 each hold one tape drive and each needs another one.

Bridge Crossing Example

- Traffic only in one direction.
- Each section of a bridge can be viewed as a resource.
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback).
- Several cars may have to be backed up if a deadlock occurs.
- Starvation is possible.

Preemptable and Nonpreemptable Resources

- A **Preemptable** resource is one that can be taken away from the process owning it with no ill effects. Such types of resources are reusable.
- Memory, printers, tape drives are example of Preemptable resources.
- A **Nonpreemptable** resource, in contrast, is one that cannot be taken away from its current owner without causing the computation to fail.
- If a process has begun to burn a CD-ROM, suddenly taking the CD recorder away from it and giving it to another process will result in a garbled CD.

Conditions for Deadlock

Deadlock can arise if four conditions hold simultaneously.

- Mutual exclusion: only one process at a time can use a resource.
- Hold and wait: a process holding at least one resource is waiting to acquire additional resources held by other processes.
- No preemption: a resource can be released only voluntarily by the process holding it, after that process has completed its task.
- Circular wait: there exists a set $\{P_0, P_1, \dots, P_{n-1}\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2 , ..., P_{n-1} is waiting for a resource that is held by P_n , and P_0 is waiting for a resource that is held by P_0 .

Methods for Handling Deadlocks

- Ensure that the system will never enter a deadlock state.
- Allow the system to enter a deadlock state and then Recover.
- Ignore the problem and pretend that deadlocks never occur in the system; used by most operating systems, including UNIX.

Deadlock Prevention

As we know, for a deadlock to occur, each of the four necessary conditions must hold. By ensuring that at least one of these conditions cannot hold, we can prevent the occurrence of a deadlock.

◆ **Mutual Exclusion**- The mutual-exclusion condition must hold for nonsharable resources. For example, a printer cannot be simultaneously shared by several processes. Sharable resources, on the other hand, do not require mutually exclusive access, and thus cannot be involved in a deadlock

◆ **Hold and Wait** – must guarantee that whenever a process requests a resource, it does not hold any other resources.

- ✦ Require process to request and be allocated all its resources before it begins execution, or allow process to request resources only when the process has none.
- ✦ Disadvantages: Low resource utilization; starvation possible.

Deadlock Prevention (Cont.)

◆ No Preemption –

- ◆ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, then all resources currently being held are released.
- ◆ Preempted resources are added to the list of resources for which the process is waiting.
- ◆ Process will be restarted only when it can regain its old resources, as well as the new ones that it is requesting.

◆ **Circular Wait** – impose a total ordering of all resource types, and require that each process requests resources in an increasing order of enumeration.

Deadlock Avoidance

Possible side effects of preventing deadlocks by Deadlock-prevention algorithms method, however, are low device utilization and reduced system throughput.

An alternative method for avoiding deadlocks is to require additional information about how resources are to be requested

- Simplest and most useful model requires that each process declare the maximum number of resources of each type that it may need.
- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
- Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes.

Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

Detection-Algorithm Usage

- ◆ When, and how often, to invoke depends on:
 - ✦ How often a deadlock is likely to occur?
 - ✦ How many processes will need to be rolled back?
 - ✓ one for each disjoint cycle
- ◆ If detection algorithm is invoked arbitrarily, there may be many cycles in the resource graph and so we would not be able to tell which of the many deadlocked processes “caused” the deadlock.

Recovery from Deadlock: Process Termination

- ◆ Abort all deadlocked processes.
- ◆ Abort one process at a time until the deadlock cycle is eliminated.
- ◆ In which order should we choose to abort?
 - ✦ Priority of the process.
 - ✦ How long process has computed, and how much longer to completion.
 - ✦ Resources the process has used.
 - ✦ Resources process needs to complete.
 - ✦ How many processes will need to be terminated.
 - ✦ Is process interactive or batch?

Recovery from Deadlock: Resource Preemption

If preemption is required to deal with deadlocks, then three issues need to be addressed:

- Selecting a victim – Which resources and which processes are to be preempted? As in process termination, we must determine the order of preemption to minimize cost.
- Rollback – return to some safe state, restart process for that state.
- Starvation – same process may always be picked as victim, include number of rollback in cost factor.

Issues Related to Deadlock:

◆ Two Phase Locking

◆ Although both avoidance and prevention are not terribly promising in the general case.

◆ As an example, in many database systems, an operation that occurs frequently is requesting locks on several records and then updating all the locked records. When multiple processes are running at the same time, there, there is a real danger of deadlock.

◆ The Approach often used is called two-phase locking.

◆ In the first phase, the process tries to lock all the records it needs, one at a time. If it succeeds, it begins the second phase, performing its updates and releasing the locks.

Issues Related to Deadlock(contd..):

◆ Non Resource Deadlocks

- ◆ This type of deadlock can occur in communication systems, in which two or more processes communicate by sending messages.
- ◆ A common arrangement is that process A sends a request message to process B, and then blocks until B sends back a reply message.
- ◆ Suppose that the request message gets lost.
- ◆ A is blocked waiting for the reply. B is blocked waiting for a request asking it to do something. We have a deadlock. Infact, there are no resources at all in the sight.

◆ Starvation

- ◆ In a dynamic system, requests for resources happen all the time. Some policy is needed to make a decision about who gets which resource when. This policy, although seemingly reasonable, may lead to some processes never getting service even though they are not deadlocked.
- ◆ This is known as starvation


Threads

- A thread, sometimes called a lightweight process (LWP), is a basic unit of CPU utilization; it comprises a thread ID, a program counter, a register set, and a stack. It is defined as the unit of dispatching
- It shares with other threads belonging to the same process its code section, data section, and other operating-system resources, such as open files and signals.
- A traditional (or heavyweight) process has a single thread of control. If the process has multiple threads of control, it can do more than one task at a time.

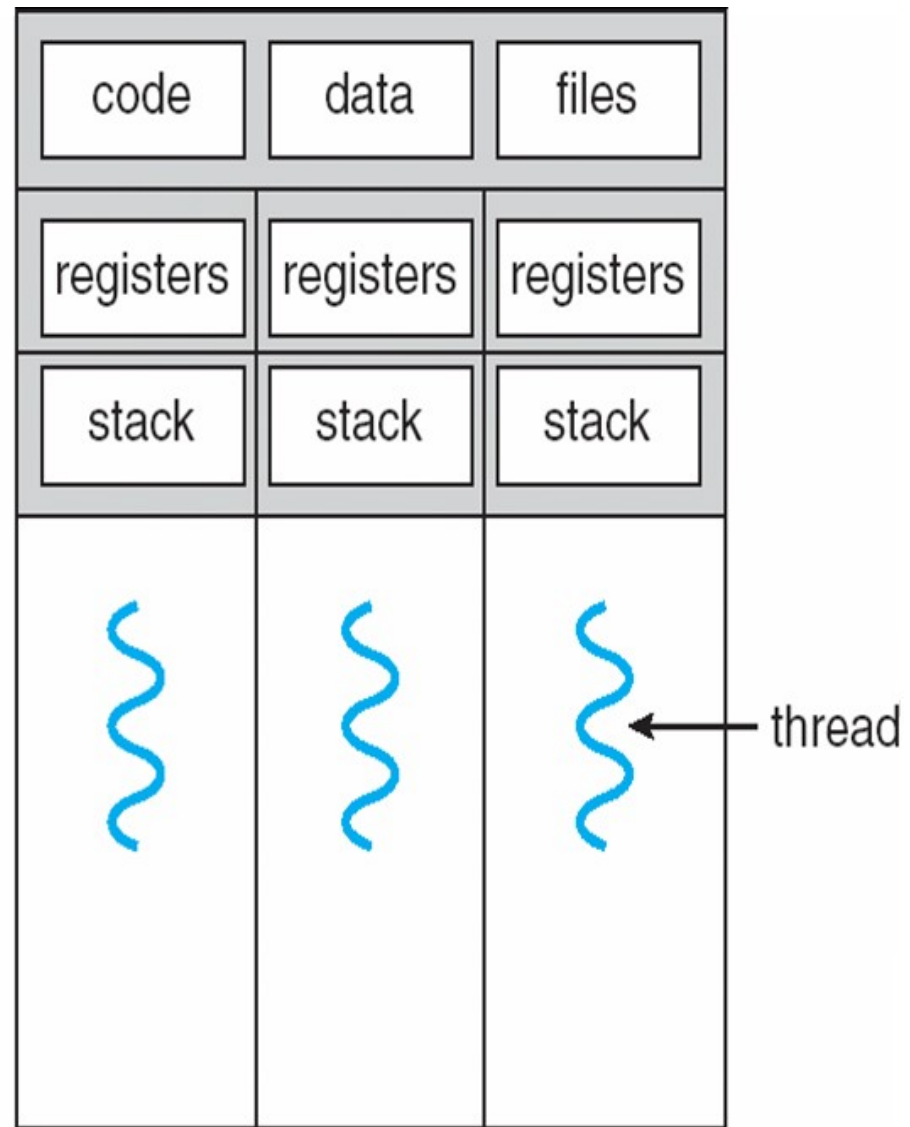
Multithreading

- Multithreading is the ability of an operating System to support multiple threads of execution within a single process
- Multiple threads run in the same address space, share the same memory areas
 - The creation of a thread only creates a new thread control structure, not a separate process image

Multithreaded Process Model




single-threaded process



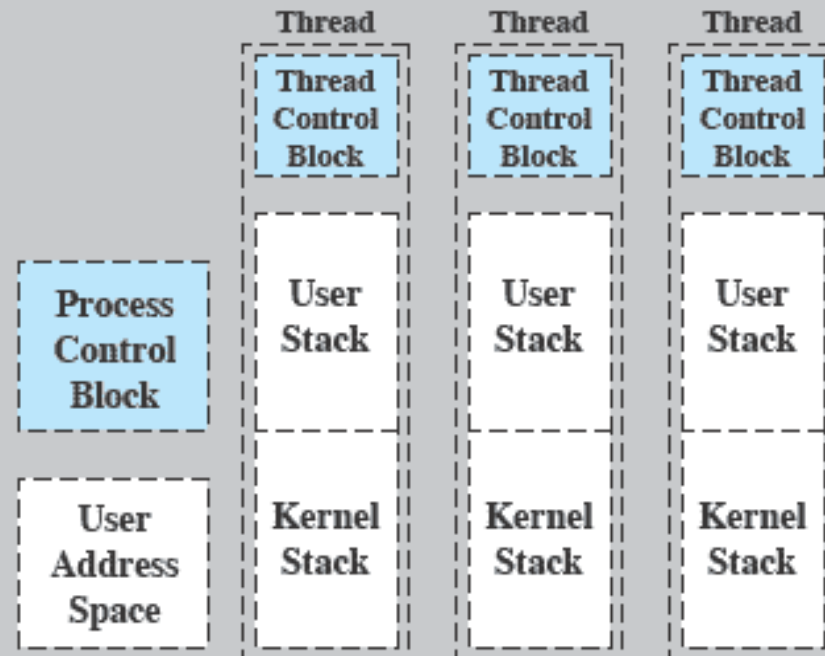
multithreaded process

Single-threaded vs Multi-threaded

Single-Threaded Process Model



Multithreaded Process Model



Thread Usage

- Many software packages that run on modern desktop PCs are multithreaded. An application typically is implemented as a separate process with several threads of control.
- A web browser might have one thread display images or text while another thread retrieves data from the network.
- A word processor may have a thread for displaying graphics, another thread for reading keystrokes from the user, and a third thread for performing spelling and grammar checking in the background.

Advantages of Threads

The benefits of multithreaded programming can be broken down into four major categories:

1. Responsiveness: Multithreading an interactive application may allow a program to continue running even if part of it is blocked or is performing a lengthy operation, thereby increasing responsiveness to the user. For instance, a multithreaded web browser could still allow user interaction in one thread while an image is being loaded in another thread.

2. Resource sharing: By default, threads share the memory and the resources of the process to which they belong. The benefit of code sharing is that it allows an application to have several different threads of activity all within the same address space.

Advantages of Threads(contd..)

3. Economy: Allocating memory and resources for process creation is costly. Alternatively, because threads share resources of the process to which they belong, it is more economical to create and context switch threads.


4. Utilization of multiprocessor architectures: The benefits of multithreading can be greatly increased in a multiprocessor architecture, where each thread may be running in parallel on a different processor. A single-threaded process can only run on one CPU, no matter how many are available. Multithreading on a multi-CPU machine increases concurrency.

User Space Threads


- User threads are supported above the kernel and are implemented by a thread library at the user level.
- The library provides support for thread creation, scheduling, and management with no support from the kernel.
- Because the kernel is unaware of user-level threads, all thread creation and scheduling are done in user space without the need for kernel intervention. Therefore, user-level threads are generally fast to create and manage.
- Its drawback is that if the kernel is single-threaded, then any user-level thread performing a blocking system call will cause the entire process to block, even if other threads are available to run within the application.

Kernel Space Threads

- Kernel threads are supported directly by the operating system. The kernel performs thread creation, scheduling, and management in kernel space.
- Because thread management is done by the operating system, kernel threads are generally slower to create and manage than are user threads.
- However, since the kernel is managing the threads, if a thread performs a blocking system call, the kernel can schedule another thread in the application for execution.
- Also, in a multiprocessor environment, the kernel can schedule threads on different processors.



User space threads




Kernel space threads

Multithreading Models

Many systems provide support for both user and kernel threads, resulting in different multithreading models.

Many-to-One Model:

- All user-level threads of one process mapped to a single kernel-level thread
- Thread management in user space
 - Efficient
 - Application can run its own scheduler implementation
- One thread can access the kernel at a time
 - Limited concurrency, limited parallelism
- Examples
 - “Green threads” (e.g. Solaris)
 - Gnu Portable Threads



Many-to-one model

Multithreading Models(contd...)

One-to-One Model:

- The one-to-one model maps each user thread to a kernel thread. It provides more concurrency than the many-to-one model by allowing another thread to run when a thread makes a blocking system call.
- It also allows multiple threads to run in parallel on multiprocessors.
- The only drawback to this model is that creating a user thread requires creating the corresponding kernel thread. Because the overhead of most implementations of this model restrict the number of threads supported by the system
- Examples:
 - Linux, along with the family of windows operating systems.

Multithreading Models(contd...)

Many-to-Many Model:

- The many-to-many model multiplexes many user-level threads to a smaller or equal number of kernel threads. The number of kernel threads may be specific to either a particular application or a particular machine.
- The many-to-many model suffers from neither of these shortcomings of above two model. Developers can create as many user threads as necessary, and the corresponding kernel threads can run in parallel on a multiprocessor.
- Also, when a thread performs a blocking system call, the kernel can schedule another thread for execution.




Fig: One-to One model




Fig: Many-to Many
model

Differences Between Processes and Threads:

- A process is a program in execution, whereas a thread is a path of execution within a process.
- Processes are generally used to execute large, 'heavyweight' jobs such as running different applications, while threads are used to carry out much smaller or 'lightweight' jobs such as auto saving a document in a program, downloading files, etc. Whenever we double-click an executable file such as Paint, for instance, the CPU starts a process and the process starts its primary thread.
- Each process runs in a separate address space in the CPU. But threads, on the other hand, may share address space with other threads within the same process. This sharing of space facilitates communication between them. Therefore, Switching between threads is much simpler and faster than switching between processes.
- Threads can directly communicate with other threads of its process; processes must use interprocess communication to communicate with sibling processes.
- Threads also have a great degree of control over other threads of the same process. Processes only have control over child processes.